




Qwen3.5-2B in llama.cpp

13th Gen Intel(R) Core(TM) i7-13620H · Q4_K_M mixed quantization · 6 CPU threads · one sequence



CPU work at 128: matrix operations are 79% of prefill and 76% of decode



MLP matrices Vocabulary head DeltaNet recurrence Vector / norm / activation
Attention / DeltaNet matrices Attention core Copies / gather / state Convolution

Top: uninstrumented medians of 1 run; whiskers show min-max. Bottom: 1 CPU trace, 8 decode steps each.
CPU operation shares do not describe Metal kernels or predict accelerator area.